

Current Research Trends

Mingyi Hong

University Of Minnesota

M. Hong would like to thank Dr. Prashant Khanduri for helping prepare the slides.

Outline

- **Communication Efficient Decentralized Optimization**
 - Motivation
 - Problem
 - Distributed: QSGD, Sparsification
 - Decentralized: Choco-Gossip based approaches
- **Optimization in the presence of Adversaries**
 - Motivation
 - Types of Byzantine Attacks
 - ByzantineSGD
 - Literature
- **Other Issues**

Communication Efficient Decentralized Optimization

Motivation

- Let us consider a distributed or decentralized architecture with m nodes
- **Goal:** To solve:

$$\min_{\mathbf{x} \in \mathbb{R}^n} f(\mathbf{x}) = \sum_{i=1}^m f_i(\mathbf{x}) \quad \text{with} \quad f_i(\mathbf{x}) = \mathbb{E}_{\xi_i} [g_i(\mathbf{x}, \xi_i)] \quad (1.1)$$

- **Typical protocol at each node**
 - Forwards an n -dimensional vector
 - Receives an n -dimensional vector
 - Update its iterate and repeat until convergence

Problem

- **Issue:** In each iteration, every node communicates an n -dimensional vector to its neighbors (or the FC)
 - Dimension n can be potentially **very large**
 - Gradients are **dense** for deep learning applications
 - Results in **network congestion** because of high communication requirements
- **Solution:** Each node communicates a **compressed** gradient
 - **Quantization:** Quantize each entry of the vector into fixed levels
 - **Sparsification:** Only send some entries (prominent ones or randomly selected) of the vector

Example

- Suppose each node uses a 32-bit arithmetic
 - **Uncompressed Communication:** Each node communicates (and receives) $32n$ bits per iteration
 - **Quantization:** 1-bit quantization
 - Only n -bits compared to $32n$ -bits
 - **Sparsification:** Sending only k out of n entries
 - $32k$ -bits instead of $32n$ -bits
- **Trade-off:** Communication vs Convergence
 - Compression \Rightarrow increased variance \Rightarrow **worse** convergence
 - Compression \Rightarrow faster communication \Rightarrow **improved** convergence

Question: Can compression lead to overall faster convergence?

Quantized SGD (QSGD)¹

- **Parallel implementation**

- Complete graph with m nodes to solve (1.1)
- **Homogeneous Data:** Each node have access to data from same distribution, i.e.,

$$f(\mathbf{x}) = f_i(\mathbf{x}) = \mathbb{E}_{\xi_i}[g_i(\mathbf{x}, \xi_i)] \quad \text{with } \xi_i \in \mathcal{D} \quad \forall i \in [m]$$

- When all the nodes have access to a common database
- Each node receives **quantized** stochastic gradient vectors from other nodes

¹Alistarh et al., QSGD: Communication-efficient SGD via gradient quantization and encoding, Advances in Neural Information Processing Systems, 2017.

Algorithm: QSGD

Quantized Stochastic Gradient Descent (QSGD)

- **For** each iteration k **do**
 - Compute stochastic gradient: $\nabla g_i(\mathbf{x}^k, \xi_i^k)$
 - Perform gradient compression: $M_i^k \leftarrow \text{Encode}(\nabla g_i(\mathbf{x}^k, \xi_i^k))$
 - Broadcast M_i^k to other nodes
 - **For** each node ℓ
 - Receive M_ℓ^k from ℓ th node
 - Decode received gradients: $\widehat{\nabla g_\ell(\mathbf{x}^k)} = \text{Decode}(M_\ell^k)$
 - **End**
 - Update iterate: $\mathbf{x}^{k+1} = \mathbf{x} - \frac{\alpha^k}{m} \sum_{\ell=1}^m \widehat{\nabla g_\ell(\mathbf{x}^k)}$
-

Same as SGD other than the Encode and Decode steps

Assumptions

Assumption 1

The function $f : \mathbb{R}^n \rightarrow \mathbb{R}$ is differentiable, convex, L -smooth, and unknown. The algorithm only has access to the stochastic gradients of f , i.e., $\nabla g(\mathbf{x}, \xi_i)$.


Assumption 2

The stochastic gradient satisfies:

- ① *Unbiased: $\mathbb{E}[\nabla g(\mathbf{x}, \xi_i)] = \nabla f(\mathbf{x})$ (We assumed $\xi_i \sim \mathcal{D} \forall i \in [m]$)*
- ② *Bounded Second Moment: **The stochastic gradient has second moment** at most B if $\mathbb{E}[\|\nabla g(\mathbf{x}, \xi_i)\|^2] \leq B$ for all $\mathbf{x} \in \mathbb{R}^n$*
- ③ *Bounded Variance: The stochastic gradient has bounded variance if $\mathbb{E}[\|\nabla f(\mathbf{x}; \xi_i) - \nabla f(\mathbf{x})\|^2] \leq \sigma^2$ for all $\mathbf{x} \in \mathbb{R}^n$*

Quantization Scheme

- **Quantization function:** $Q_s(v)$
 - $s \geq 1$ uniformly distributed levels between 0 and 1
 - Unbiased quantization and introduces minimal variance



Quantization with $s = 5$ levels

- **Efficient Coding of Gradients** $Q_s(v) = (\|v\|, \sigma, s \cdot \zeta)$
 - $\|v\|$: 32-bit representation
 - σ : Sign of individual elements: 1-bit representation
 - $s \cdot \zeta$: Integer quantization levels: Use Elias codes²

²Elias, Universal codeword sets and representations of the integers, IEEE Transactions on Information Theory, 1975.

Stochastic Quantization

- For any $\mathbf{v} \in \mathbb{R}^n$ with $\mathbf{v} \neq 0$, $Q_s(\mathbf{v})$ is defined as

$$Q_s(v_i) = \|\mathbf{v}\| \cdot \text{sgn}(v_i) \cdot \zeta_i(\mathbf{v}, s) \quad \text{for } i \in [n]$$

where $\text{sgn}(v_i)$ is the sign function and $\zeta_i(\mathbf{v}, s)$'s are independent random variables defined as:

- Let $0 \leq \ell < s$ be s.t. $|v_i|/\|\mathbf{v}\| \in [\ell/s, (\ell+1)/s]$, then

$$\zeta_i(\mathbf{v}, s) = \begin{cases} \ell/s & \text{w.p. } 1 - p\left(\frac{|v_i|}{\|\mathbf{v}\|}, s\right) \\ (\ell+1)/s & \text{otherwise} \end{cases}$$

with $p(a, s) = as - \ell$ for any $a \in [0, 1]$. $Q(\mathbf{v}, s) = 0$ if $\mathbf{v} = 0$

Distribution of $\zeta_i(\mathbf{v}, s)$ has minimal variance over distributions with support $\{0, 1/s, \dots, 1\}$ and $\mathbb{E}[\zeta_i(\mathbf{v}, s)] = |v_i|/\|\mathbf{v}\|$

Main Result

Theorem 1.1

- Under Assumptions 1 and 2, and $\alpha^k = 1/(L + \sqrt{m}/\gamma)$ and to achieve $\mathbb{E}\left[f\left(\frac{1}{K} \sum_{k=1}^K \mathbf{x}^k\right)\right] - f^* \leq \epsilon$, QSGD requires K iterations with

$$K = O\left(R^2 \cdot \max\left\{\frac{2B'}{m\epsilon^2}, \frac{L}{\epsilon}\right\}\right).$$

with $\gamma = \frac{1}{\sigma} \sqrt{\frac{2}{K}}$, $R^2 = \sup_{\mathbf{x} \in \mathbb{R}^n} \|\mathbf{x} - \mathbf{x}^0\|$, $\sigma = B'$,
 $B' = \min\{n/s^2, \sqrt{n}/s\}B$

- Moreover, if we take $s = \sqrt{n}$, **QSGD** requires $2.8n + 32$ bits per iteration compared to $32n$ required by **SGD**

Comparison to SGD

Network	Dataset	Params.	Init. Rate	Top-1 (32bit)	Top-1 (QSGD)	Speedup (8 GPUs)
AlexNet	ImageNet	62M	0.07	59.50%	60.05% (4bit)	2.05 ×
ResNet152	ImageNet	60M	1	77.0%	76.74% (8bit)	1.56 ×
ResNet50	ImageNet	25M	1	74.68%	74.76% (4bit)	1.26 ×
ResNet110	CIFAR-10	1M	0.1	93.86%	94.19% (4bit)	1.10 ×
BN-Inception	ImageNet	11M	3.6	-	-	1.16 × (projected)
VGG19	ImageNet	143M	0.1	-	-	2.25 × (projected)
LSTM	AN4	13M	0.5	81.13%	81.15 % (4bit)	2 × (2 GPUs)

Speed-up achieved by **QSGD** compared to **SGD** while training different networks with the percentages indicating the percentage of time consumed in communication process

SignSGD³

- **SignSGD:** 1-bit compression
 - Transmitting just the **sign** of each stochastic gradient
- **SignSGD** works well when
 - Gradients are as **dense**
 - Both gradients and noise are dense in deep learning problem
- **Homogeneous Data:** Each node have access to data from same distribution, i.e.,

$$f(\mathbf{x}) = f_i(\mathbf{x}) = \mathbb{E}_{\xi_i}[g_i(\mathbf{x}, \xi_i)] \text{ with } \xi_i \sim \mathcal{D} \forall i \in [m]$$

- **Non-Convex** functions

³Bernstein et al., SignSGD: Compressed Optimisation for Non-Convex Problems, Thirty-fifth International Conference on Machine Learning, 2018

Algorithm: SignSGD

SignSGD

- **Input** Learning rate α , initial iterate \mathbf{x}^0
 - **For** $k = 0$ to K **do**
 - Compute Stochastic Gradient: $\nabla g(\mathbf{x}^k, \xi^k)$ with $|\xi^k| = b_k$
 - Update: $\mathbf{x}^{k+1} = \mathbf{x}^k - \alpha \operatorname{sign}(g(\mathbf{x}^k, \xi^k))$
 - **End For**
-

Algorithm: SignSGD with Majority Vote

Distributed SignSGD with Majority Voting (MV)

- **Input** Learning rate α , initial iterate \mathbf{x}^0 , m nodes each with gradient estimate $\nabla g_i(\mathbf{x}^k, \xi_i^k)$ for $i \in [m]$ and with $|\xi_i^k| = b_k$
- **For** $k = 0$ to $K - 1$ **do**
- **on** Central Node
- **pull** $\text{sign}(\nabla g_i(\mathbf{x}^k, \xi_i^k))$ **from** each node
- **push** $\text{sign}[\sum_{i=1}^m \text{sign}(\nabla g_i(\mathbf{x}^k, \xi_i^k))]$ **to** each node (MV)
- **on** each worker update:

$$\mathbf{x}^{k+1} = \mathbf{x}^k - \alpha \text{sign} \left[\sum_{i=1}^m \text{sign}(\nabla g_i(\mathbf{x}^k, \xi_i^k)) \right]$$

- **End For**
-

Assumptions

Assumption 3 (Smooth)

The function f is non-convex and for all $\mathbf{x}, \mathbf{y} \in \mathbb{R}^n$ we have

$$\|\nabla f_i(\mathbf{y}) - \nabla f_i(\mathbf{x})\| \leq L_i \|\mathbf{y} - \mathbf{x}\|$$

for a vector of non-negative constants $\vec{L} = [L_1, \dots, L_n]$

Assumption 4 (Variance Bound)

The stochastic gradient computed at each $j \in [m]$ for each dimension $i \in [n]$ satisfies

$$\mathbb{E}[\|\nabla g_j(\mathbf{x}, \xi_j)_i - \nabla g_j(\mathbf{x})_i\|^2] \leq \sigma_i^2$$

for a vector of non-negative constants $\vec{\sigma} = [\sigma_1, \dots, \sigma_n]$

Main Results: SignSGD

Theorem 1.2

For iterations 1 to K under Assumptions 3 and 4, with $\alpha = 1/\sqrt{\|\vec{L}\|_1 K}$ and mini-batch size $b_k = K$. Let N be the cumulative number of stochastic gradient calls up to step K , i.e., $N = O(K^2)$. Then for **SignSGD** we have

$$\mathbb{E} \left[\frac{1}{K} \sum_{k=0}^{K-1} \|\nabla f(\mathbf{x}^k)\|_1 \right]^2 \leq \frac{1}{\sqrt{N}} \left[\sqrt{\|\vec{L}\|_1} \left(f_0 - f^* + \frac{1}{2} \right) + 2\|\vec{\sigma}\|_1 \right]^2$$

Note that the norms are ℓ_1 -norms!

Main Results: SignSGD with Majority Voting

Theorem 1.3

- **SignSGD with Majority Vote** with m workers converges at least as fast as **SignSGD**
- Further Assuming that the noise in each component of the stochastic gradient is unimodal and symmetric about the mean we have for **SignSGD with Majority Vote**

$$\begin{aligned} \mathbb{E} \left[\frac{1}{K} \sum_{k=0}^{K-1} \|\nabla f(\mathbf{x}^k)\|_1 \right]^2 \\ \leq \frac{1}{\sqrt{N}} \left[\sqrt{\|\vec{L}\|_1} \left(f_0 - f^* + \frac{1}{2} \right) + \frac{2}{\sqrt{m}} \|\vec{\sigma}\|_1 \right]^2 \end{aligned}$$

Comparison of SignSGD to SGD

- **Define:**

- Density of a vector: $\phi(\mathbf{v}) = \frac{\|\mathbf{v}\|_1^2}{n\|\mathbf{v}\|_2^2}$
- Lower bound on gradient density: $\phi(\nabla g)$
- $R_1 = \frac{\sqrt{\phi(\vec{L})}}{\phi(\nabla g)}$ and $R_2 = \frac{\phi(\vec{\sigma})}{\phi(\nabla g)}$ with

- Convergence of **SignSGD** can be rephrased as:

$$\mathbb{E} \left[\frac{1}{K} \sum_{k=0}^{K-1} \|\nabla f(\mathbf{x}^k)\|_2 \right]^2 \leq \frac{2}{\sqrt{N}} \left[R_1 L \left(f_0 - f^* + \frac{1}{2} \right)^2 + 4R_2 \sigma^2 \right]$$

with $\sigma^2 = \|\vec{\sigma}\|_2^2$ and $L = \|\vec{L}\|_\infty$

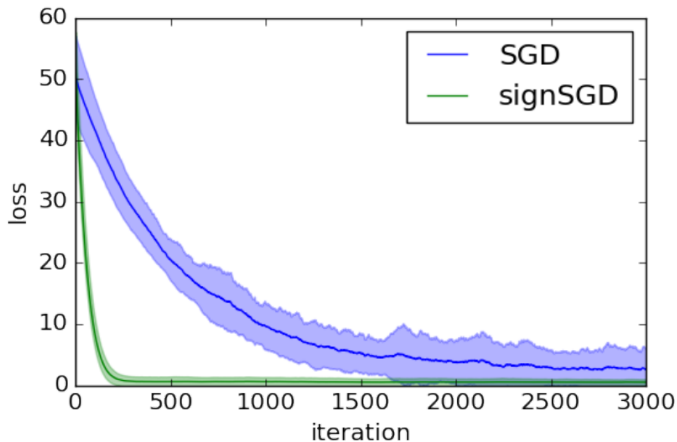
- For **SGD**

$$\mathbb{E} \left[\frac{1}{K} \sum_{k=0}^{K-1} \|\nabla f(\mathbf{x}^k)\|_2^2 \right] \leq \frac{1}{\sqrt{N}} [2L(f_0 - f^*) + \sigma^2]$$

Comparison of SignSGD to SGD

- $R_1 \gg 1$ and $R_2 \gg 1$
 - **SGD** is better suited than **SignSGD**
 - **Curvature and the stochasticity are much denser than the typical gradient**
- NOT[$R_1 \gg 1$] and NOT[$R_2 \gg 1$]
 - **SignSGD** is as fast or faster than **SGD**
 - **Neither curvature nor stochasticity are much denser than the gradient**

SGD vs SignSGD



Toy example with $f(x) = \frac{1}{2}\|x\|_2^2$ for $x \in \mathbb{R}^{100}$.

Literature

- **Lossy compression:** Low precision arithmetic⁴
- **Sparsification**⁵
- **Three levels**⁶

⁴Abadi et al., Tensorflow: Large-scale machine learning on heterogeneous distributed systems, arXiv 2016.

⁵De Sa, Christopher M., et al. "Taming the wild: A unified analysis of hogwild-style algorithms." Advances in neural information processing systems. 2015.

⁶Wen et al., Terngrad: Ternary gradients to reduce communication in distributed deep learning, NeurIPS 2017.

Compression for Decentralized Optimization: Choco-Gossip⁷

- **Decentralized** networks
- General compression
 - Quantize or Sparsify, Biased or unbiased
- **Choco-Gossip**: Decentralized consensus with **compression**
 - Achieves **exact consensus** while converging to the true solution
- **Past Approaches**: Only **neighborhood convergence** or required **fairly accurate compression**
- **Choco-SGD**: Decentralized SGD based on **Choco-Gossip**

⁷Koloskova et al., Decentralized Stochastic Optimization and Gossip Algorithms with Compressed Communication, 36th International Conference on Machine Learning, 2019.

Model

- **Goal:** Solve (1.1) over a decentralized network
- **Network:** Undirected network $\mathcal{G} = \{\mathcal{V}, \mathcal{E}\}$ with $|\mathcal{V}| = m$ nodes and $|\mathcal{E}|$ edges
 - Matrix $W = [w_{ij}] \in [0, 1]^{m \times m}$ is symmetric doubly stochastic
 - Denote by $\delta > 0$ the spectral gap of W
 - Denoting $\beta = \|I - W\|_2 \in [0, 2]$
- **Heterogeneous Data:** Each node have access to data from potentially different distributions, i.e.,

$$f_i(\mathbf{x}) = \mathbb{E}_{\xi_i}[g_i(\mathbf{x}, \xi_i)] \text{ with } \xi_i \sim \mathcal{D}_i \ \forall i \in [m]$$

Compression

A general notion of compression operator Q :

Assumption 5 (Compression Operator)

We assume the compression operator $Q : \mathbb{R}^n \rightarrow \mathbb{R}^n$ satisfies for all $\mathbf{x} \in \mathbb{R}^n$

$$\mathbb{E}_Q \|Q(\mathbf{x}) - \mathbf{x}\|^2 \leq (1 - \omega) \|\mathbf{x}\|^2$$

for $\omega > 0$. Here, \mathbb{E}_Q is the expectation over the randomness of Q

- **QSGD** discussed earlier satisfies Assumption 5
- **Sparsification** satisfies Assumption 5

Choco-SGD: Based on Choco-Gossip

Algorithm: Choco-SGD

- **Initialize:** Initial iterate $\mathbf{x}_i^0 \in \mathbb{R}^n \ \forall i \in [m]$, consensus step-size γ , SGD step-size α^k , Graph \mathcal{G} , W , initialize $\hat{\mathbf{x}}_i^0 = 0 \ \forall i \in [m]$
 - **For** $k = 0$ to $K - 1$ **do**
 - Compute stochastic gradient $\mathbf{d}_i^k = \nabla g_i(\mathbf{x}_i^k, \xi_i^k)$
 - $\mathbf{x}_i^{k+\frac{1}{2}} = \mathbf{x}_i^k - \alpha^k \mathbf{d}_i^k$
 - $\mathbf{q}_i^k = Q(\mathbf{x}_i^{k+\frac{1}{2}} - \hat{\mathbf{x}}_i^k)$
 - Send \mathbf{q}_i^k and receive \mathbf{q}_j^k from neighbours
 - Maintain: $\hat{\mathbf{x}}_j^{k+1} = \mathbf{q}_j^k + \hat{\mathbf{x}}_j^k$ for all $j \in \mathcal{N}_i$
 - $\mathbf{x}_i^{k+1} = \mathbf{x}_i^{k+\frac{1}{2}} + \gamma \sum_{j:\{i,j\} \in \mathcal{E}} w_{ij} (\hat{\mathbf{x}}_j^{k+1} - \hat{\mathbf{x}}_i^{k+1})$
 - **End For**
-

Intuition: Choco-SGD

- Note that at every iteration $k \in [K]$ each node $i \in [m]$
 - Maintains a **compressed estimate** of \mathbf{x}_j^k denoted as:

$$\hat{\mathbf{x}}_j^k \text{ for } j : \{i, j\} \in \mathcal{E} \text{ (including } \{i\} \in \mathcal{E})$$

- This is accomplished by receiving **compressed error estimate**

$$\mathbf{q}_j^k = Q(\mathbf{x}_i^{k+\frac{1}{2}} - \hat{\mathbf{x}}_i^k) \text{ for } j : \{i, j\} \in \mathcal{E} \text{ (including } \{i\} \in \mathcal{E})$$

at each $k \in [K]$

- **Idea behind Choco-SGD:** Note that averaging the update equation across all $i \in [m]$, yields the standard SGD update

$$\bar{\mathbf{x}}^{k+1} = \bar{\mathbf{x}}^{k+\frac{1}{2}} = \bar{\mathbf{x}}^k - \alpha^k \bar{\mathbf{d}}^k$$

where $\bar{\mathbf{x}}^k = \frac{1}{m} \sum_{i=1}^m \mathbf{x}_i^k$ and $\bar{\mathbf{d}}^k$ is defined in a similar fashion

Assumptions

Assumption 6

Each $f_i : \mathbb{R}^n \rightarrow \mathbb{R}$ for $i \in [m]$ is assumed to be L -smooth and μ -strongly convex and variance of each node $i \in [m]$ is bounded for all $\mathbf{x} \in \mathbb{R}^n$ as

$$\begin{aligned}\mathbb{E}_{\xi_i} \|\nabla f_i(\mathbf{x}, \xi_i) - \nabla f_i(\mathbf{x})\|^2 &\leq \sigma_i^2 \\ \mathbb{E}_{\xi_i} \|\nabla f_i(\mathbf{x}, \xi_i)\|^2 &\leq G^2\end{aligned}$$

where \mathbb{E}_{ξ_i} denotes expectation over $\xi_i \sim \mathcal{D}_i$. Moreover, we denote

$$\bar{\sigma}^2 = \frac{1}{m} \sum_{i=1}^m \sigma_i^2$$

Main Result

Theorem 1.4

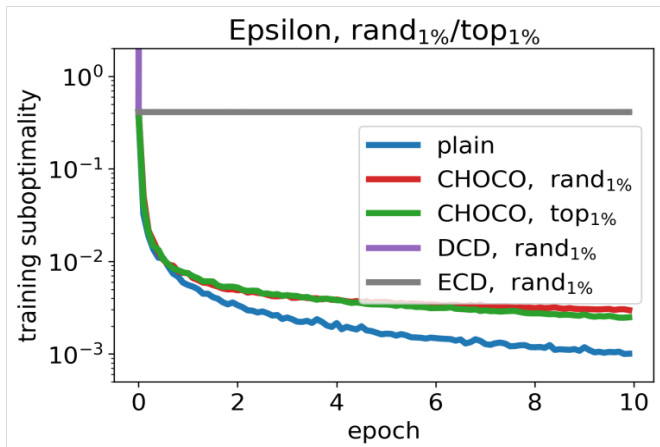
Under Assumptions 6 and 5 the algorithm **Choco-SGD** with step-sizes $\alpha^k = \frac{4}{\mu(a+k)}$ with $a \geq \max\{410/\delta^2\omega, 16\kappa\}$ for $\kappa = L/\mu$ and $\gamma = \frac{\delta^2\omega}{16\delta+\delta^2+4\beta^2+2\delta\beta^2-8\delta\omega}$ converges with a rate

$$f(\mathbf{x}_{\text{avg}}^K - f^*) = O\left(\frac{\bar{\sigma}^2}{\mu m K}\right) + O\left(\frac{\kappa G^2}{\mu \omega^2 \delta^4 K^2}\right) + O\left(\frac{G^2}{\mu \omega^3 \delta^6 K^3}\right)$$

where $\mathbf{x}_{\text{avg}}^K = \frac{1}{S_K} \sum_{k=0}^{K-1} w_k \bar{\mathbf{x}}^k$ and $S_K = \sum_{k=0}^{K-1} w_k$

For large K , the term $O\left(\frac{\bar{\sigma}^2}{\mu m K}\right)$ dominates recovering the rate of SGD

Comparison: Choco-SGD vs SGD, ECD and DCD⁸



Comparison of **Choco-SGD** with Plain **SGD** with **ECD-SGD** and **DCD-SGD** proposed in [Tang et al. 2018] with $\text{rand}_{1\%}$ sparsification and $\text{top}_{1\%}$ for Choco-SGD for RCV1 dataset

⁸Tang et al., Decentralized training over decentralized data, ICML 2018

Extensions

- **Non-Convex functions**⁹
- **Push-Sum + Choco-Gossip**¹⁰
- **Gradient Descent for Strongly Convex function**¹¹

⁹Singh, et al., SPARQ-SGD: Event-Triggered and Compressed Communication in Decentralized Stochastic Optimization." arXiv 2019.

¹⁰Taheri, et al., Quantized Decentralized Stochastic Learning over Directed Graphs." arXiv 2020.

¹¹Liu et al., Linear Convergent Decentralized Optimization with Compression." arXiv 2020.

Optimization in the Presence of Adversaries

The Byzantine Generals Problem¹²

- Adversaries also referred to as **Byzantines**
- Byzantine army communicating only by messages
- One or more generals may be traitors who will try to confuse the other generals

Goal: The algorithm must ensure that the generals agree upon a common battle plan

¹²Lamport et al., "The Byzantine Generals Problem", ACM 1982

Optimization in Presence of Byzantines


- **Byzantine nodes:** A fraction of multiple WNs might be adversarial
 - Forward **arbitrary** vectors instead of gradients to the server
 - Can collaborate and adversely effect the algorithm's performance

Problem: To design optimization algorithms which are resilient to Byzantine attacks

- **Next:** Types of Byzantine attacks

Classic Byzantine Attacks

- A **fraction of WNs are Byzantines**¹³
 - The set of Byzantine nodes is fixed
- A **fraction of WNs are Byzantines**, however, the set of Byzantine nodes is **not fixed**¹⁴
 - Nodes can change alliances




Classical Byzantine Attacks

¹³Alistarh et al., Byzantine Stochastic Gradient Descent, NeurIPS, 2018.


¹⁴Blanchard et al., Machine Learning with Adversaries: Byzantine Tolerant Gradient Descent, NeurIPS, 2017.

Generalized Byzantine Attacks

- **Dimensional Byzantine Attacks**¹⁵
 - The attacker only affects certain **dimensions** of a node's vector



(a) Classic Byzantine



(b) Generalized Byzantine

Dimension based attacks

¹⁵Xie et al., Generalized Byzantine-tolerant SGD, arxiv, 2018.

Byzantine SGD¹⁶

- **Problem:** Solve (1.1) in a distributed fashion
- **Model:**
 - Number of nodes: m out of which β -fraction are Byzantines
 - Nodes interact via a server node (**star network**)
 - Each node forwards their stochastic gradient to the server and receives a aggregated direction
 - Server **filters** the alleged Byzantine nodes
 - Classical Byzantine attack model
- **Goal:** The algorithm must ensure convergence in the presence of Byzantines!
 - By designing effective filtering rule

¹⁶Alistarh et al., Byzantine Stochastic Gradient Descent, NeurIPS, 2018.

Notations

- **Homogeneous Data:** Each node have access to data from same distribution, i.e.,

$$f(\mathbf{x}) = f_i(\mathbf{x}) = \mathbb{E}_{\xi_i}[g_i(\mathbf{x}, \xi_i)] \text{ with } \xi_i \sim \mathcal{D} \ \forall i \in [m]$$

- The set of good nodes: \mathcal{G}
 - \mathcal{G} is not known to the server
- **Estimate** of the good set at each iteration: \mathcal{G}^k for $k \in [K]$

Assumptions

Assumption 7

For differentiable $f : \mathbb{R}^n \rightarrow \mathbb{R}$ we have

- f is μ -strongly convex
- f is L -Lipschitz smooth
- f is B -Lipschitz continuous (Gradient of f is bounded)

At each iteration, each node $i \in [m]$ computes $\nabla_i^k \in \mathbb{R}^n$ as

Assumption 8

At each iteration $k \in [K]$ for every $i \in \mathcal{G}$, we have

- $\nabla_i^k = \nabla g_i(\mathbf{x}^k, \xi_i^k)$ for a random sample $\xi_i^k \sim \mathcal{D}$
- $\|\nabla_i^k - \nabla f(\mathbf{x}^k)\| \leq \mathcal{V}$

For each $k \in [K]$ and $i \notin \mathcal{G}$ the vector ∇_i^k can be adversarially chosen

Filtering Rule

- Let us first define two **statistics** the central node maintains for filtering the Byzantine nodes
 - $A_i = \sum_{t=1}^k \langle \nabla_i^k, \mathbf{x}^t - \mathbf{x}^1 \rangle$
 - $B_i = \sum_{t=1}^k \nabla_i^t$
- Note that A_i and B_i accumulate over time
 - The filtering rule relies on the **Martingale concentration**
 - Relies on the fact the for $i \in \mathcal{G}$ and $k \in [K]$ the stochastic gradients ∇_i^k are **chosen independently**
- **Vector Median:** Finally, we define vector median of a set of vectors $\mathbf{v}_1, \dots, \mathbf{v}_m$ as any vector \mathbf{v}_i such that

$$|\{j \in [m] : \|\mathbf{v}_j - \mathbf{v}_i\| \leq \mathfrak{T}_v\}| > m/2$$

where $\mathfrak{T}_v > 0$ is the diameter of the norm-ball

Algorithm: Byzantine SGD

Input: Learning rate α , Initial iterate \mathbf{x}^1 , constants $\mathfrak{T}_A, \mathfrak{T}_B > 0$

- $\mathcal{G}^1 \leftarrow [m]$
- **For** $k = 1$ to K **do**
- Receive $\nabla_i^k \in \mathbb{R}^n$ from $i \in [m]$
- Maintain statistics A_i and B_i
- Compute $A_{\text{med}} = \text{median}\{A_1, \dots, A_m\}$
- Compute B_{med} from B_i 's with diameter \mathfrak{T}_B (see previous slide)
- Compute ∇_{med} from ∇_i^k 's with diameter $2\mathcal{V}$ (see previous slide)
- Filtering Rule:

$$\mathcal{G}^k \leftarrow \{i \in \mathcal{G}^{k-1} : |A_i - A_{\text{med}}| \leq \mathfrak{T}_A \cap \|B_i - B_{\text{med}}\| \leq \mathfrak{T}_B \\ \cap \|\nabla_i^k - \nabla_{\text{med}}\| \leq 4\mathcal{V}\}$$

- $\mathbf{x}^{k+1} = \underset{\mathbf{y}: \|\mathbf{y} - \mathbf{x}^1\| \leq D}{\text{argmin}} \left\{ \frac{1}{2} \|\mathbf{y} - \mathbf{x}^k\|^2 + \alpha \left\langle \frac{1}{m} \sum_{i \in \mathcal{G}^k} \nabla_i^k, \mathbf{y} - \mathbf{x}^k \right\rangle \right\}$
- **End For**

Main Result

Theorem 1.5

Under Assumption 7 with $\mu = 0$ (convex function) and with β -fraction of Byzantine nodes with $\beta < 1/2$, the ByzantineSGD finds a point x with $f(x) - f^ \leq \epsilon$ in K iterations where*

$$K = \tilde{O}\left(\frac{1}{\epsilon} + \frac{1}{\epsilon^2 m} + \frac{\beta^2}{\epsilon^2}\right)$$

or

$$K = \tilde{O}\left(\frac{1}{\mu} + \frac{1}{\mu \epsilon m} + \frac{\beta^2}{\mu \epsilon}\right) \text{ if } f(x) \text{ is } \mu\text{-strongly convex}$$

Note that $\beta = 0$ leads to the guarantees of standard parallel SGD!

Past: Deterministic Optimization Approaches

- **Byzantine Gradient Descent**¹⁷
 - Classical Byzantine model, Strongly convex functions, Median based aggregation
- **Approximate Gradient Descent**¹⁸
 - Classical Byzantine model, Strongly convex functions, Direction with most spread
- **Robust Distributed Gradient Descent**¹⁹
 - Generalized Byzantine model, Strongly convex, convex and non-convex functions, Coordinate wise median and trimmed mean

¹⁷Chen et al., Distributed Statistical Machine Learning in Adversarial Settings: Byzantine Gradient Descent, Proc. ACM 2017.

¹⁸Su et al., Securing Distributed Machine Learning in High Dimensions, arxiv 2018.

¹⁹Yin et al., Byzantine-Robust Distributed Learning: Towards Optimal Statistical Rates, ICML 2018.

Past: Stochastic Gradient Descent

- **Krum Based Approaches**²⁰

- Classical Byzantine model, Non-convex functions, Geometric median based aggregation

- **Phocas**²¹

- Generalized Byzantine model, Non-convex and strongly convex functions, Coordinate wise trimmed mean

- **Robust Gradient Aggregation (RSA)**²²

- Classical Byzantine model, Strongly convex functions, Model based aggregation, **Heterogeneous** data

²⁰Blanchard et al., Machine Learning with Adversaries: Byzantine Tolerant Gradient Descent, NeurIPS, 2017

²¹Xie et al., "Phocas: dimensional Byzantine-resilient stochastic gradient descent", arxiv 2018

²²Li et al., RSA: Byzantine-Robust Stochastic Aggregation Methods for Distributed Learning from Heterogeneous Datasets, AAAI, 2019.

Past: Decentralized Gradient Decent

- **BRIDGE**²³ and **ByRDIE**²⁴
 - Generalized Byzantine models, Strongly convex and convex functions resp., Distance based aggregation
- **Recent Survey**²⁵

²³Yang et al., BRIDGE: Byzantine-resilient decentralized gradient descent, arXiv 2019.

²⁴Yang et al., ByRDIE: Byzantine-resilient distributed coordinate descent for decentralized learning, IEEE Transactions on Signal and Information Processing over Networks 2019.

²⁵Yang et al., Adversary-resilient distributed and decentralized statistical inference and machine learning: An overview of recent advances under the Byzantine threat model, IEEE Signal Processing Magazine 2020.