

Problem Set 1. Solution

2020-11

1 Problem 1

- (1) Show that $X \sim \mathcal{N}(0, 1)$ is the maximum entropy distribution such that $\mathbb{E}X = 0$ and $\mathbb{E}X^2 = 1$.

Solution.

$$\begin{aligned} \min_p \quad & \int_{\mathcal{X}} p(x) \log p(x) \, dx \\ \text{s.t.} \quad & \int_{\mathcal{X}} p(x) \, dx = 1 \\ & \int_{\mathcal{X}} x p(x) \, dx = 0 \\ & \int_{\mathcal{X}} x^2 p(x) \, dx = 1 \end{aligned}$$

The Lagrangian is

$$\mathcal{L}(p; \lambda, \mu) = \int p(x) \log p(x) \, dx + \lambda_0 \left(\int p(x) \, dx - 1 \right) + \lambda_1 \int x p(x) \, dx + \lambda_2 \left(\int x^2 p(x) \, dx - 1 \right)$$

which is convex in p . Then taking

$$\frac{\partial \mathcal{L}}{\partial p} = \log p(x) + 1 + \lambda_0 + \lambda_1 x + \lambda_2 x^2 = 0$$

we have

$$p(x) = \exp \left(-(\lambda_0 + \lambda_1 x + \lambda_2 x^2) - 1 \right) \geq 0.$$

$$p(x) = \exp \left(-(\lambda_0 + \lambda_1 x + \lambda_2 x^2) - 1 \right) \text{ with } \int x p(x) \, dx = 0, \text{ we have } \lambda_1 = 0.$$

$$p(x) = \exp \left(-(\lambda_0 + \lambda_2 x^2) - 1 \right) \text{ with } \int x^2 p(x) \, dx = 1 \text{ and } \int p(x) \, dx = 1, \text{ we have } \lambda_0 = \log \sqrt{2\pi} - 1, \lambda_2 = \frac{1}{2}.$$

$$\text{Therefore, } p(x) = \frac{1}{\sqrt{2\pi}} \exp \left(-\frac{1}{2}x^2 \right), \text{ i.e., } X \sim \mathcal{N}(0, 1). \quad \blacksquare$$

- (2) Generalize the result in (1) for the maximum entropy distribution given the first k moments, i.e., $\mathbb{E}X^i = m_i, i = 1, \dots, k$.

Solution. Write the problem as

$$\begin{aligned} \min_p \quad & \int_{\mathcal{X}} p(x) \log p(x) \, dx \\ \text{s.t.} \quad & \int_{\mathcal{X}} x^n p(x) \, dx = m_n, \quad n = 0, \dots, k, \quad m_0 := 1 \end{aligned}$$

The Lagrangian is

$$\mathcal{L}(p; \lambda) = \int p(x) \log p(x) \, dx + \sum_{0 \leq n \leq k} \lambda_n \left(\int x^n p(x) \, dx - m_n \right)$$

which is convex in p . Then taking

$$\frac{\partial \mathcal{L}}{\partial p} = \log p(x) + 1 + \sum_{0 \leq n \leq k} \lambda_n x^n = 0$$

we have

$$p(x) = \exp \left(-\sum_{0 \leq n \leq k} \lambda_n x^n - 1 \right),$$

an exponential family, then λ is determined by the constraints $\int x^n p(x) \, dx = m_n, n = 0, \dots, k. \quad \blacksquare$

2 Problem 2

Let Y_1, \dots, Y_n be a set of independent random variables with the following pdfs

$$p(y_i | \theta_i) = \exp(y_i b(\theta_i) + c(\theta_i) + d(y_i)), \quad i = 1, \dots, n$$

Let $\mathbb{E}(Y_i) = \mu_i(\theta_i)$, $g(\mu_i) = x_i^\top \beta$, where g is the link function and $\beta \in \mathbb{R}^d$ is the vector of model parameters.

(1) Denote $g(\mu_i)$ as η_i , and let s be the score function of β . Show that

$$s_j = \sum_{i=1}^n \frac{(y_i - \mu_i) x_{ij}}{\text{Var}(Y_i)} \frac{\partial \mu_i}{\partial \eta_i}, \quad j = 1, \dots, d$$

Solution. Let $L_i := \log p(y_i | \theta_i) = y_i b(\theta_i) + c(\theta_i) + d(y_i)$, then

$$s_j = \sum_{i=1}^n \frac{\partial L_i}{\partial \beta_j} = \sum_{i=1}^n \frac{\partial L_i}{\partial \theta_i} \frac{\partial \theta_i}{\partial \mu_i} \frac{\partial \mu_i}{\partial \eta_i} \frac{\partial \eta_i}{\partial \beta_j} = \sum_{i=1}^n \left(y_i \frac{\partial b(\theta_i)}{\partial \theta_i} + \frac{\partial c(\theta_i)}{\partial \theta_i} \right) \frac{\partial \theta_i}{\partial \mu_i} \frac{\partial \mu_i}{\partial \eta_i} x_{ij}, \quad j = 1, \dots, d.$$

By interchanging the differentiation $\frac{\partial}{\partial \theta_i}$ and integration $\mathbb{E}1, \mathbb{E}Y_i$, we have

$$\begin{aligned} 0 &= \frac{\partial}{\partial \theta_i} \mathbb{E}1 = \mu_i \frac{\partial b(\theta_i)}{\partial \theta_i} + \frac{\partial c(\theta_i)}{\partial \theta_i} \\ \frac{\partial \mu_i}{\partial \theta_i} &= \frac{\partial}{\partial \theta_i} \mathbb{E}Y_i = (\mathbb{E}Y_i^2) \frac{\partial b(\theta_i)}{\partial \theta_i} + \mu_i \frac{\partial c(\theta_i)}{\partial \theta_i}, \end{aligned}$$

then

$$\begin{aligned} y_i \frac{\partial b(\theta_i)}{\partial \theta_i} + \frac{\partial c(\theta_i)}{\partial \theta_i} &= \left(y_i \frac{\partial b(\theta_i)}{\partial \theta_i} + \frac{\partial c(\theta_i)}{\partial \theta_i} \right) - \left(\mu_i \frac{\partial b(\theta_i)}{\partial \theta_i} + \frac{\partial c(\theta_i)}{\partial \theta_i} \right) = (y_i - \mu_i) \frac{\partial b(\theta_i)}{\partial \theta_i} \\ \frac{\partial \mu_i}{\partial \theta_i} &= \frac{\partial \mu_i}{\partial \theta_i} - \mu_i \left(\mu_i \frac{\partial b(\theta_i)}{\partial \theta_i} + \frac{\partial c(\theta_i)}{\partial \theta_i} \right) = \text{Var}(Y_i) \frac{\partial b(\theta_i)}{\partial \theta_i}. \end{aligned}$$

therefore

$$s_j = \sum_{i=1}^n \frac{(y_i - \mu_i) \frac{\partial b(\theta_i)}{\partial \theta_i}}{\text{Var}(Y_i) \frac{\partial b(\theta_i)}{\partial \theta_i}} \frac{\partial \mu_i}{\partial \eta_i} x_{ij} = \sum_{i=1}^n \frac{(y_i - \mu_i) x_{ij}}{\text{Var}(Y_i)} \frac{\partial \mu_i}{\partial \eta_i}, \quad j = 1, \dots, d.$$

■

(2) Let \mathcal{I} be the Fisher information matrix. Show that

$$\mathcal{I}_{jk} = \mathbb{E}(s_j s_k) = \sum_{i=1}^n \frac{x_{ij} x_{ik}}{\text{Var}(Y_i)} \left(\frac{\partial \mu_i}{\partial \eta_i} \right)^2, \quad \forall 1 \leq j, k \leq d$$

Solution.

$$\begin{aligned} \mathcal{I}_{jk} &= \mathbb{E}(s_j s_k) \\ &= \mathbb{E} \left(\sum_{i=1}^n \frac{(y_i - \mu_i) x_{ij}}{\text{Var}(Y_i)} \frac{\partial \mu_i}{\partial \eta_i} \sum_{l=1}^n \frac{(y_l - \mu_l) x_{lk}}{\text{Var}(Y_l)} \frac{\partial \mu_l}{\partial \eta_l} \right) \\ &= \sum_{1 \leq i, l \leq n} \frac{x_{ij} x_{lk} \mathbb{E}(y_i - \mu_i)(y_l - \mu_l)}{\text{Var}(Y_i) \text{Var}(Y_l)} \frac{\partial \mu_i}{\partial \eta_i} \frac{\partial \mu_l}{\partial \eta_l} \\ &= \sum_{i=1}^n \frac{x_{ij} x_{ik}}{\text{Var}(Y_i)} \left(\frac{\partial \mu_i}{\partial \eta_i} \right)^2, \quad \text{by independence.} \end{aligned}$$

■

3 Problem 3

Use the following code to generate co-variate matrices X .

```
1 import numpy as np
2
3 np.random.seed(1234)
4 n = 100
5 X = np.random.normal(size=(n,2))
```

- (1). Generate $n = 100$ observations Y following the logistic regression model with true parameter $\beta_0 = (-2, 1)$.

Solution.

```
1 import numpy as np
2 from scipy.special import expit as sigmoid
3
4 def generate_data(beta0, n=100):
5     X = np.random.normal(size=(n, beta0.shape[0]))
6     logits = X @ beta0
7     probs = sigmoid(logits)
8     return {"X":X, "logits":logits, "probs":probs}
9
10 def generate_y(probs):
11     y = np.random.binomial(1, probs)
12     return y
13
14 seed = 1234
15 np.random.seed(seed)
16 beta0 = np.array([-2., 1.])
17 data = generate_data(beta0, n=100)
18 X, probs = data["X"], data["probs"]
19 y = generate_y(probs=probs)
20 print(y)

```

```
1 [0 0 1 0 0 1 0 0 0 0 1 0 0 0 1 1 0 0 0 0 1 0 0 1 1 0 1 1 0 1 1 1 0 0 1 1 0
2  1 0 0 0 0 1 1 0 0 1 0 0 0 1 0 0 0 0 0 1 0 0 1 0 1 1 1 1 1 0 0 0 1 1 0 0 0 1
3  1 1 0 0 1 1 1 0 1 0 0 0 1 0 0 0 1 0 1 1 1 0 0 0 0 0]
```

- (2). Find the MLE using the iteratively reweighted least square algorithm.

Solution.

```
1 import numpy as np
2 from scipy import sparse
3 from scipy.special import expit as sigmoid
4 from numpy.linalg import inv as inverse
5 from numpy.linalg import norm
6
```

```

7 def irls_lr(X, y, max_itr=200, epsilon=1e-12, quiet=True):
8     ''' Readme:
9         this function is defined to estimate
10        the parameters of logistic regression model
11        by IRLS, Iteratively Reweighted Least Square Algorithm.
12    '''
13    # initialization
14    n, d = X.shape
15    beta = inverse(X.T@X) @ (X.T@y)
16    W = sparse.dia_matrix((n, n))
17    err_path = []
18
19    # main
20    for i in range(1, max_itr+1):
21        logits = X @ beta
22        probs = sigmoid(logits)
23        W = sparse.diags(probs*(1-probs))
24        beta_ = beta + inverse(X.T@W@X) @ (X.T@y-probs)
25        err = norm(beta_-beta)/norm(beta)
26        err_path.append(err)
27        beta = beta_
28
29        if i % 5 == 0:
30            if not quiet:
31                print(f"err: {err:.3e}, itr: {i}")
32        if err < epsilon:
33            if not quiet:
34                print(f"err: {err:.3e}, itr: {i}")
35                print(f"MLE: {list(beta)}")
36            break
37    # returns
38    out = {"beta":beta, "itr":i, "errs":err_path}
39    return out

```

```

1 out = irls_lr(X, y, quiet=0)

```

```

1 err: 2.217e-06, itr: 5

```

```

2 err: 1.627e-16, itr: 7

```

```

3 MLE: [-1.3708659534317205, 0.6698777671314895]

```

- (3). Repeat (1) and (2) for 100 instances. Compare the MLEs with the asymptotical distribution $\hat{\beta} \sim \mathcal{N}(\beta_0, \mathcal{I}^{-1}(\beta_0))$. Present your result with a scatter plot for MLEs with contours for the PDF of the asymptotical distribution.

Solution.

```

1 W = sparse.diags(probs*(1-probs))
2 Fisher = X.T @ W @ X
3 cov = inverse(Fisher)
4 print(cov)

```

```

1 [[ 0.19350715 -0.0440933 ]
2  [-0.0440933  0.10205911]]

```

```

1 def experiment(
2     X, probs, num_tri=100,
3     generate_y=generate_y, alg=irls_lr,
4 ):
5     instances = [generate_y(probs) for i in range(1, num_tri+1)]
6     betas = [alg(X, y)["beta"] for y in instances]
7     return np.array(betas)
8
9 np.random.seed(1234)
10 beta0 = np.array([-2., 1.])
11 betas = experiment(X, probs, num_tri=100)
12 print(np.cov(betas.T))

```

```

1 [[ 0.23124879 -0.06162206]
2  [-0.06162206  0.10223682]]

```

```

1 from scipy import stats
2 import matplotlib.pyplot as plt
3
4 radius = 0.8
5 x0 = np.linspace(beta0[0]-radius, beta0[0]+radius, 50)
6 x1 = np.linspace(beta0[1]-radius, beta0[1]+radius, 50)
7 X0, X1 = np.meshgrid(x0, x1) # create x0-x1 meshgrid
8 pos = np.dstack((X0, X1))     # shape=(num1,num2,d)
9
10 binorm_rv1 = stats.multivariate_normal(beta0, cov)
11 Z1 = binorm_rv1.pdf(pos)
12
13 plt.figure(figsize=(6, 6))
14 plt.plot(betas[:, 0], betas[:, 1], 'o', mfc='none') # scatter
15 CS = plt.contour(X0, X1, Z1, cmap='cool')          # contour
16 plt.clabel(CS, inline=1)
17 plt.title(r'Comparison ($n=100$)')
18 plt.savefig('./comparison_100.pdf', bbox_inches='tight')
19 plt.show()

```

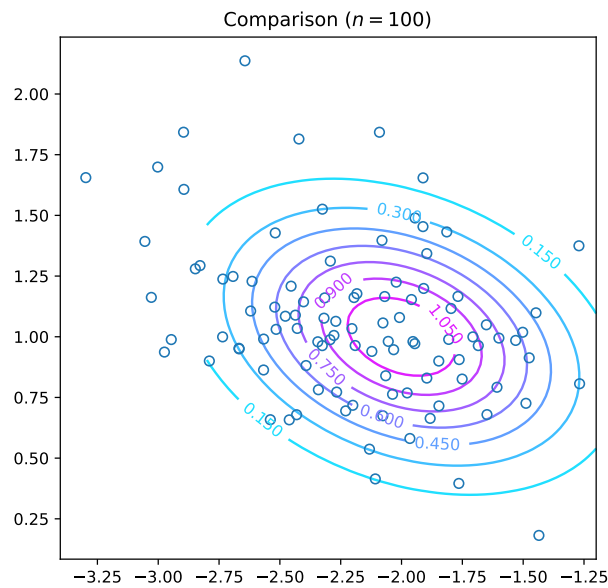


Figure 3.1: Comparison

- (4). Try the same for $n = 10000$. Does the asymptotical distribution provide a better fit to the MLEs? You can use the empirical covariance matrix of the MLEs for comparison.

Solution.

```

1 seed = 1234
2 np.random.seed(seed)
3 beta0 = np.array([-2., 1.])
4
5 data = generate_data(beta0, n=100)
6 X, probs = data["X"], data["probs"]
7 betas_ = experiment(X, probs, num_tri=100)
8 betas_cov_ = np.cov(betas_.T)
9
10 W = sparse.diags(probs*(1-probs))
11 Fisher = X.T @ W @ X
12 cov_ = inverse(Fisher)
13 print(cov_)
14 print(betas_cov_)
15 err = np.linalg.norm(cov_-betas_cov_, "fro")
16 print(f"err:{err:.6f}, err_rel:{err/np.linalg.norm(cov_, 'fro'):.6f}")
17
18 data = generate_data(beta0, n=10000)
19 X, probs = data["X"], data["probs"]
20 betas = experiment(X, probs, num_tri=100)
21 betas_mean = np.mean(betas, axis=0)
22 betas_cov = np.cov(betas.T)
23 W = sparse.diags(probs*(1-probs))
24 Fisher = X.T @ W @ X

```

```

25 cov = inverse(Fisher)
26 print(betas_cov)
27 print(cov)
28 err = np.linalg.norm(cov-betas_cov,"fro")
29 print(f"err:{err:.6f}, err_rel:{err/np.linalg.norm(cov,'fro'):.6f}")

```

```

1 [[ 0.19350715 -0.0440933 ]
2  [-0.0440933  0.10205911]]
3 [[ 0.25623927 -0.04276872]
4  [-0.04276872  0.12578528]]
5 err:0.067095, err_rel:0.294943
6 [[ 0.00202756 -0.00069806]
7  [-0.00069806  0.00105249]]
8 [[ 0.00174037 -0.00053367]
9  [-0.00053367  0.00096422]]
10 err:0.000380, err_rel:0.178520

```

```

1 plt.figure(figsize=(12, 6))
2
3 radius = 0.4
4 plt.subplot(121)
5 x0 = np.linspace(beta0[0]-radius, beta0[0]+radius, 50)
6 x1 = np.linspace(beta0[1]-radius, beta0[1]+radius, 50)
7 X0, X1 = np.meshgrid(x0, x1) # create x0-x1 meshgrid
8 pos = np.dstack((X0, X1))     # shape-(num1,num2,d)
9
10 binorm_rv1 = stats.multivariate_normal(beta0, cov_)
11 Z1 = binorm_rv1.pdf(pos)
12 plt.plot(betas[:, 0], betas[:, 1], 'o', mfc='none') # scatter
13 CS1 = plt.contour(X0, X1, Z1, cmap='cool')         # contour
14 plt.clabel(CS1, inline=1)
15 plt.title(r'Comparison ($n=10000$)')
16 plt.savefig('./comparison_10000.pdf', bbox_inches='tight')
17
18 plt.subplot(122)
19 x2 = np.linspace(betas_mean[0]-radius, betas_mean[0]+radius, 50)
20 x3 = np.linspace(betas_mean[1]-radius, betas_mean[1]+radius, 50)
21 X2, X3 = np.meshgrid(x2, x3) # create x0-x1 meshgrid
22 pos2 = np.dstack((X2, X3))   # shape-(num1,num2,d)
23
24 binorm_rv2 = stats.multivariate_normal(betas_mean, betas_cov)
25 Z2 = binorm_rv2.pdf(pos2)
26 CS1 = plt.contour(X0, X1, Z1, cmap='cool')
27 plt.clabel(CS1, inline=1)
28 CS2 = plt.contour(X2, X3, Z2, cmap='spring')

```

```

29 plt.xlabel(CS2, inline=1)
30 plt.title(r'Comparison with asymptotical distribution')
31 plt.savefig('./Asymptotical distribution.pdf', bbox_inches='tight')
32
33 plt.show()

```

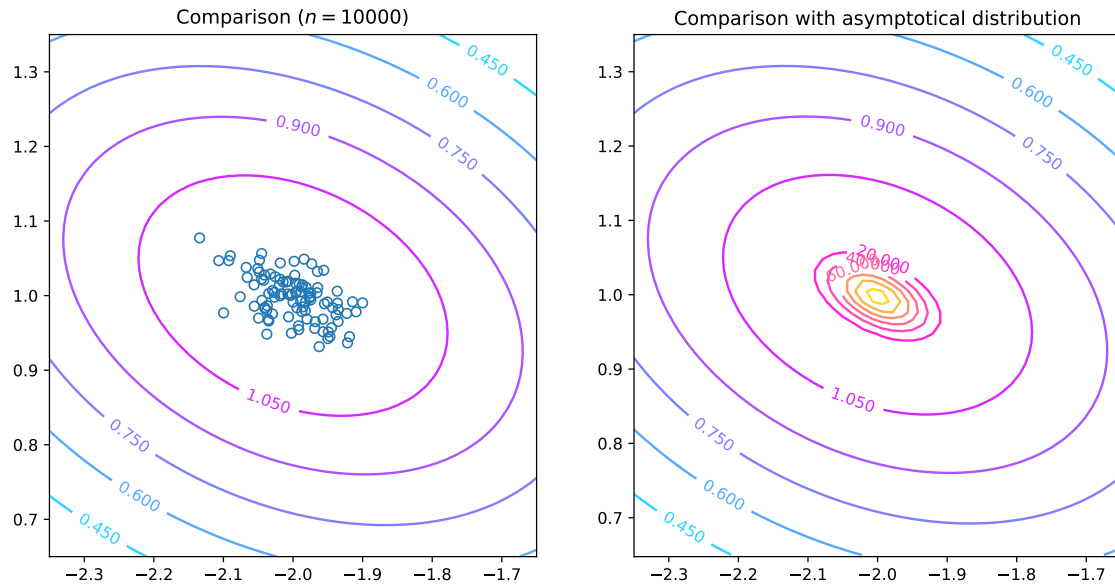


Figure 3.2: Comparison

The asymptotic distribution with the larger sample size, which has a more intense concentration and a lower error rate, suggests a better fit for the MLEs.

4 Problem 4

Consider the probit regression model

$$Y | X, \beta \sim \text{Bernoulli}(p), \quad p = \Phi(X\beta)$$

where Φ is the cumulative distribution function of the standard normal distribution. Similarly as in Problem 3, generate a large covariate matrix X with 100000 instances and 100 features, and response Y with true parameter β_0

```

1 import numpy as np
2 np.random.seed(1234)
3
4 n, d = 100000, 100
5 X = np.random.normal(size=(n, d))
6 beta_0 = np.random.normal(size=d)

```

(1). Compare gradient descent and Nesterov's accelerated gradient descent.

Solution.

```

1 def GD(self, init, tol = 1e-6, step_size = 0.0001, maxit = 1000):

```



```

2     beta_old = np.zeros(shape = (np.shape(init)[0],)) + 1
3     beta_new = init
4     likelihood = np.array([])
5     l = 1000
6     l_next = 1
7     for i in range(1, maxit + 1):
8         l = l_next
9         if i % 100 == 0 :
10            print(l)
11            grad = self.gradient(beta_new)
12            beta_old = beta_new
13            beta_new = beta_old + step_size * grad
14            l_next = self.loglikelihood(beta_new)
15            likelihood = np.append(likelihood, l_next)
16            if abs(l - l_next) / abs(l) < tol:
17                break
18            return beta_new, likelihood
19
20 def NAG(self, init, tol = 1e-6, step_size = 0.0001, maxit = 1000):
21     beta_previous = init
22     beta_now = init
23     beta_new = init
24     likelihood = np.array([])
25     l = 1000
26     l_next = 1
27     for i in range(1, maxit + 1):
28         l = l_next
29         beta_previous = beta_now
30         beta_now = beta_new
31         if i % 100 == 0 :
32             print(l)
33         y = beta_now + (i - 2) / (i + 1) * (beta_now - beta_previous)
34         grad = self.gradient(y)
35         beta_new = y + step_size * grad
36         l_next = self.loglikelihood(beta_new)
37         likelihood = np.append(likelihood, l_next)
38         if abs(l - l_next) / abs(l) < tol:
39             break
40     return beta_new, likelihood

```




Figure 4.3: Comparison: SGD and NAG

Although the objective function is non-monotone decreasing when using nesterov's acceleration, it can reach a much optimal value in a fairly less amount of time.

- (2). Compare vanilla stochastic gradient descent with different adaptive stochastic gradient descent methods, including AdaGrad, RMSprop, and Adam. Using minibatch sizes 32,64,128.

Solution.

```

1 def SGD(self, init, tol = 1e-6, step_size = 0.1, batch_size = 32, maxit = 2500):
2     beta_old = np.zeros(shape = (np.shape(init)[0],)) + 1
3     beta_new = init
4     likelihood = np.array([])
5     l = 1000
6     l_next = 1
7     for i in range(1, maxit + 1):
8         l = l_next
9         if i % 100 == 0 :
10             print(l)
11         grad = self.stoch_grad(beta_new, batch = batch_size)
12         beta_old = beta_new
13         beta_new = beta_old + step_size * grad
14         l_next = self.loglikelihood(beta_new)
15         likelihood = np.append(likelihood, l_next)
16         if abs(l - l_next) / abs(l) < tol:
17             break
18     return beta_new, likelihood
19
20 def AdaGrad(self, init, tol = 1e-10, epsilon = 1e-8, step_size = 0.2, batch_size = 32,
21             maxit = 2500):
22     beta_new = init
23     likelihood = np.array([])
24     l = 1000
25     l_next = 1
26     grad_sum = 0
27     for i in range(1, maxit + 1):
28         l = l_next

```

```

29         if i % 100 == 0 :
30             print(l)
31             grad = self.stoch_grad(beta_new, batch = batch_size)
32             grad_sum = grad_sum + np.square(grad)
33             beta_new = beta_new + step_size * grad / np.sqrt(grad_sum + epsilon)
34             l_next = self.loglikelihood(beta_new)
35             likelihood = np.append(likelihood, l_next)
36             if abs(l - l_next) / abs(l) < tol:
37                 break
38         return beta_new, likelihood
39
40 def RMSprop(self, init, tol = 1e-10, epsilon = 1e-8, step_size = 0.01, batch_size = 32,
41             maxit = 2500):
42     beta_new = init
43     likelihood = np.array([])
44     l = 1000
45     l_next = 1
46     grad_sum = 0
47     for i in range(1, maxit + 1):
48         l = l_next
49         if i % 100 == 0 :
50             print(l)
51             grad = self.stoch_grad(beta_new, batch = batch_size)
52             grad_sum = 0.9 * grad_sum + 0.1 * np.square(grad)
53             beta_new = beta_new + step_size * grad / np.sqrt(grad_sum + epsilon)
54             l_next = self.loglikelihood(beta_new)
55             likelihood = np.append(likelihood, l_next)
56             if abs(l - l_next) / abs(l) < tol:
57                 break
58     return beta_new, likelihood
59
60 def Adam(self, init, tol = 1e-10, epsilon = 1e-8, step_size = 0.008, batch_size = 32,
61          maxit = 2500):
62     beta1 = 0.9
63     beta2 = 0.999
64     beta_new = init
65     likelihood = np.array([])
66     l = 1000
67     l_next = 1
68     grad_sum = 0
69     grad_square = 0
70     for i in range(1, maxit + 1):
71         l = l_next
72         if i % 100 == 0 :
73             print(l)
74             grad = self.stoch_grad(beta_new, batch = batch_size)
75             grad_sum = beta1 * grad_sum + (1 - beta1) * grad

```

```

76     grad_square = beta2 * grad_square + (1 - beta2) * np.square(grad)
77     beta1_power = beta1 ** (i + 1)
78     beta2_power = beta2 ** (i + 1)
79     mt = grad_sum / (1 - beta1_power)
80     vt = grad_square / (1 - beta2_power)
81     beta_new = beta_new + step_size * mt / np.sqrt(vt + epsilon)
82     l_next = self.loglikelihood(beta_new)
83     likelihood = np.append(likelihood, l_next)
84     if abs(l - l_next) / abs(l) < tol:
85         break
86     return beta_new, likelihood

```




Figure 4.4: Comparison: batch size = 32, 64, 128

We can see that the adaptive learning rate methods can rapidly fall to a minimum value at the beginning of training, and this becomes more obvious when we enlarge batch size. Among the four, AdaGrad performs the best, while RMSProp fluctuates a lot. As the batch size enlarges, all of the adaptive learning rate methods can reach a lower value.

- (3). Bonus question. Generate a random mask matrix M as follows and use it to sparsify the covariance matrix X

```

1  np.random.seed(1234)
2
3  sparse_rate = 0.3
4  M = np.random.uniform(size=(n,d)) < sparse_rate
5  X[M] = 0.

```

Repeat your experiments in (2), and compare with the results for the full covariance matrix

Solution.




Figure 4.5: Comparison: batch size = 32, 64, 128

In this setting, the difference between SGD and its adaptive variants becomes more obvious, especially in the case of large batch size. Other observations are roughly the same as (2).